An introduction to MatPIV v. 1.6.1

Article · May 2011

1 author:

Johan Kristian Sveen
Institute for Energy Technology
24 PUBLICATIONS   **1,140** CITATIONS

An introduction to **MatPIV** v. 1.6.1

J. Kristian Sveen

August 6, 2004

Abstract

Particle Image Velocimetry (PIV) has seen a rapid growth over the last two decades, much owing to the developments in digital camera and solid state laser technologies. PIV can essentially be looked upon as an application of pattern matching principles to experiments. We rely upon hardware such as lasers and cameras for illumination and image capture. Subsequently we use computer code to perform the pattern matching. **MatPIV** is one of a variety of different computer codes available written specifically for this purpose. The vast majority of codes are commercially available, but in the last 7-8 years several Open Source PIV codes have been created and are currently being more or less actively maintained. MatPIV is one of these codes and it is distributed under the GNU General Public License[1]

This document acts as the entry level tutorial for using the MatPIV code. Very basic theory is reviewed and references to appropriate sources are included. The focus, however, is on the use of **MatPIV** and how it is implemented.

---

[1]**MatPIV** is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version. This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

The GNU General Public License can be found on the World Wide Web at http://www.gnu.org/copyleft/gpl.html or it can be obtained by writing to the Free Software Foundation, Inc., 59 Temple Place - Suite 330, Boston, MA 02111-1307, USA.

Contents

Introduction

MatPIV is a program written for the authors personal educational purposes only and should be treated thereafter. This document acts primarily as an entry level documentation to the MatPIV m-files and the Particle Image Velocimetry method used here. Users/Readers are advised to consult the review paper on PIV [Sveen and Cowen, 2004], a decent book on the subject, e.g. Raffel et al. [1998], the PhD thesis by Westerweel [1993], and finally the various articles written by different authors over the last decade or so (some references are included in the bibliography list).

The **MatPIV** distribution comes with 3 different sets of demo-images which are taken from the papers by Grue et al. [1999] and Jensen et al. [2001].

The **MatPIV** files have been tested on Linux machines only, with MATLAB versions from 5.2.1.1420 and up to 6.5.0. It is recommended that users have the Image processing toolbox, although a few workarounds do exist.

The aim of the present document is primarily to give an overview of the functionality of the **MatPIV** toolbox. At the time of writing, **MatPIV** is one of at least three available, free toolboxes. It is, however, by far the largest presently available, both when functionality and number of users are considered. Interested readers may google for **UraPIV** and **mpiv** to find alternative PIV-software.

Chapter 1

Installation

MatPIV is distributed as zipped archives which may be downloaded from the **MatPIV** webpage.

1.1 Unix/Linux

This comes in the form of a single tar.gz file which should be installed in the following steps: 1) create a directory called, say, MatPIV1.6.1 and cd to it. 2) type

```
~> tar -zxvf MatPIV1.6.1.tar.gz
```

1.2 Windows

This comes in the form of a zip-archive. Install by first creating a directory called, say, MatPIV1.6.1 and then unzipping the archive into it (using your preferred zip-utility).

1.3 Adding the path to Matlab

Subsequently you need to add this directory and all it's subdirectories to the **Matlab** path. This can be done in a few different ways. Firstly, if you are running **Matlab** with it's full desktop, you can click 'File', 'Set Path', followed by 'Add with sub-folders' and finally adding your newly created directory.

Alternatively for Unix/Linux users, you can add a file named *startup.m* to your home-directory, containing the following lines

```
p=genpath('/the/path/where/you/installed/MatPIV1.6.1');
addpath(p);
```

This assumes you always start up **Matlab** from the root of your home-directory. If not, the paths are not added.

Chapter 2

An example session

In this section we will take a look at an example session with **MatPIV**. The images can be found in the "Demo3"-subdirectory of your **MatPIV** -installation, and they are taken from an experiment concerning surface waves in water (Jensen et al. [2001]). The images are shown in figure 2.1.

2.1   The easy way

We can quite easily perform our first PIV-calculation by changing to the *Demo3* directory and executing the command

```
>> [x,y,u,v]=matpiv('mpim1b.bmp','mpim1c.bmp',64,0.0012,0.5,'single');
```

This will interrogate the images in a single pass using 64×64 pixels large interrogation windows with 50% overlap. The term "single pass" refers to the fact that the calculation is performed by looping over the images with only one iteration. It is normally better to use more iterations and achieve higher accuracy, but we shall return to these subjects at a later stage.

The result consists of four matrices *x, y, u* and *v* which are measured in pixels and pixels/second. The result can be visualized by issuing the following command:

```
>> quiver(x,y,u,v)
```

This is all well, but normally we prefer to know velocities and positions in engineering units such as centimeters and seconds. To do this we need to know how large each pixel in our image is, measured in physical coordinates. Furthermore, it may in some cases be nice to mask out particular regions of the flow, in order to save computational time. For example, there is no need to perform our calculations in the airy part above the wave-crest in figure 2.1.

2.2   The not so easy way

We will now go through the complete process of measuring the velocities from the accompanying demo-images. The process is in our case started by first defining a coordinate system so that we know how far the particles move, measured in physical coordinates (centimeters or meters) instead of only "camera" coordinates (pixels). Thereafter we apply a mask to the images to avoid performing calculations where there are no particles (outside the flow field). Subsequently the images will be interrogated using a window-shifting technique [Westerweel et al., 1997]. After the velocity field has been found, a series of filters are applied before finally all the identified outliers are interpolated using a nearest neighbor interpolation.

2.2.1   Specifying the coordinate system

In this tutorial we will focus on the practical aspects of PIV and it's application to real experiments. Our first step to quantify the velocity fields from the images is to determine how large each pixel is in the

Figure 2.1: Surface wave. Field-of-view (FOV) is roughly $30 * 30$ cm

images. We can do this by inserting a grid with points into the field of view and take an image of it with
our camera. Such an image is shown in figure 2.2, which shows 10 points that are 5cm apart. We can
now use the function *definewoco.m* to calculate how large a pixel is in the image and the orientation of our
coordinate system. In this example we will use the 6 "dots" in the image to define our coordinate system.
In this specific case, the dots are rather large and have a flat peak. Therefore we will use a special option in
definewoco that cross-correlates the image with a gaussian bell to emphasize the center of the peaks. The
peaks in question are about 8 pixels in diameter, so we'll pass this information to definewoco.

```
>> definewoco('mpwoco.bmp','.');
Please give the approximate width of your points (in pixels -
default is 20). Type 0 here to get old behaviour of definewoco: 8
....calculating....this may take a few seconds.
...Done!
Now mark the crosses you whish to use as coordinate points
Please mark your world coordinate points with left mouse button.
Press ENTER when finished!
Now you need to give the physical coordinates to each of the points specified!
-----------------------
Please enter the world coordinates for the white
 circle  marked in the current figure (in square parenthesis): [-10 10]
Please enter the world coordinates for the white
 circle  marked in the current figure (in square parenthesis): [-10 5]
Please enter the world coordinates for the white
 circle  marked in the current figure (in square parenthesis): [0 10]
Please enter the world coordinates for the white
 circle  marked in the current figure (in square parenthesis): [0 5]
Please enter the world coordinates for the white
 circle  marked in the current figure (in square parenthesis): [10 5]
Please enter the world coordinates for the white
 circle  marked in the current figure (in square parenthesis): [10 10]
Mapping function. (N)onlinear or (L)inear (N/n/L/l): l
Error (norm) = 0.019268
Save coordinate file....specify number >>
Coordinate mapping factors saved to file:  worldco
```

Figure 2.2: Image of coordinate system.

In another set of images (Demo 1) included with **MatPIV** (where the coordinate image file is called *woco.bmp*) the peaks are well defined and we would use the more standard approach

```
>> definewoco('woco.bmp','o');
```

2.2.2 Masking out regions of the flow

We observe that there is an area above the water surface where there are no particles. In this region we do not want to perform our PIV calculations and hence we can use the file *mask.m*[1] to mask out this region of the flow. The file displays the image and the mask is defined by clicking with the left mouse button, and defining the final point using the middle mouse button. Figure 2.3

```
>> mask('mpim1b.bmp','worldco.mat');
Mark your polygon points with the left mouse button.
Press the middle button when you are finished, press
<BACKSPACE> or <DELETE> to remove the previously selected vertex.
Do you whish to add another field to mask? (1 = yes, 0 = no) >> 0
* Calculating the pixel to world transformation using linear mapping - DONE
```

2.2.3 Calculating velocities

Now we would like to calculate the velocities from the images. We use the window shifting technique mentioned above and start off with 64×64 large images and end with 16×16 after 6 iterations. The images are taken with a time separation of 0.0012s and we will use 50% overlap of the interrogation regions.

---

[1]The file *mask.m* uses a function from the Image Processing Toolbox by the Mathworks Inc. For those that do not possess this toolbox there exists a workaround version called *mask2.m*

Figure 2.3: Masked image. Red region will be excluded in the subsequent calculations.

```
>> [x,y,u,v,snr,pkh]=matpiv('mpim1b.bmp','mpim1c.bmp',...
[64 64;64 64;32 32;32 32;16 16;16 16],...
0.0012,0.5,'multin','worldco.mat','polymask.mat');

* Pass No: 1
 No. of vectors: 1094 , Seconds taken: 19.606562.
 Global filter running - with limit: 3 *std [U V] ..... 2 vectors changed
 Local median filter running: ..........................
............2 vectors changed.
 Interpolating outliers: ....4 Nan's interpolated.
   Expanding velocity-field for next pass
 Interpolating outliers: .1 Nan's interpolated.
* Pass No: 2
 No. of vectors: 1094 , Seconds taken: 20.628889.
 Global filter running - with limit: 3 *std [U V] ..... 0 vectors changed
 Local median filter running: ..........................
............3 vectors changed.
 Interpolating outliers: ...3 Nan's interpolated.
   Expanding velocity-field for next pass
 Interpolating outliers: ...............218 Nan's interpolated.
* Pass No: 3
 No. of vectors: 4402 , Seconds taken: 17.721477.
 Global filter running - with limit: 3 *std [U V] ..... 1 vectors changed
 Local median filter running: ..........................
..............................................
.....3 vectors changed.
```

```
 Interpolating outliers: ....4 Nan's interpolated.
   Expanding velocity-field for next pass
 Interpolating outliers: .1 Nan's interpolated.
* Pass No: 4
 No. of vectors: 4402 , Seconds taken: 17.555643.
 Global filter running - with limit: 3 *std [U V] ..... 0 vectors changed
 Local median filter running: ...........................
...............................................
.....1 vectors changed.
 Interpolating outliers: .1 Nan's interpolated.
   Expanding velocity-field for next pass
 Interpolating outliers: ............................460 Nan's interpolated.
* Pass No: 5
 No. of vectors: 17673 , Seconds taken: 26.574317.
 Global filter running - with limit: 3 *std [U V] ..... 24 vectors changed
 Local median filter running: ...........................
..........................................
..................................................
........................257 vectors changed
.
 Interpolating outliers: ................................
...............281 Nan's interpolated.
* Final Pass
 - Using 16*16 interrogation windows!
 No. of vectors: 17673, Seconds taken: 33.649626
* Calculating the pixel to world transformation using linear mapping - DONE
```

2.2.4 Filtering the result

The next step is to filter the velocity-fields to remove so called spurious vectors. These are vectors that occur primarily due to low image quality in some parts of the image(s). We may, for example, have regions with insufficient seeding (not enough particles to create a good pattern for matching), or too many particles so that the image is saturated locally. Most users should find that applying some or all the following should work well. In the final step we replace the missing vector values using a nearest neighbor interpolation.

```
>> [su,sv]=snrfilt(x,y,u,v,snr,1.3);
 SnR filter running with threshold value = 1.3  - finished...
391 outliers identified

>> [pu,pv]=peakfilt(x,y,su,sv,pkh,0.5);
 Peak height filter running .......
329 vectors changed

>> [gu,gv]=globfilt(x,y,pu,pv,3);
 Global filter running - with limit: 3 *std [U V] .....
1 vectors changed

>> [mu,mv]=localfilt(x,y,gu,gv,2,'median',3,'polymask.mat');
 Local median filter running: ...............................
.........................................................
.........................................................
.........1464 vectors changed.
```

Figure 2.4: Filtered and interpolated velocity-field.


Figure 2.5: Magnitude of velocity.

```
>> [fu,fv]=naninterp(mu,mv,'linear','polymask.mat',x,y);
 Interpolating outliers: ...................................
...................2161 Nan's interpolated.
```

2.2.5   Visualizing the results

Figure 2.4 shows $1/4$ of the velocity-vectors plotted using the **Matlab** command *quiver*. In this case only every fourth vector is shown due to the high number of vectors present. The figure was created with the following commands:

```
>> quiver(x(1:4:end),y(1:4:end),fu(1:4:end),fv(1:4:end),2); axis tight
```

It is also possible to use the function *magnitude.m* to calculate the magnitude of each velocity vector, $(\text{fu}^2 + \text{fv}^2)^{1/2}$. Figure 2.5 was created using the following commands:

```
>> w=magnitude(x,y,fu,fv);
>> pcolor(x,y,w), shading flat, colorbar
```

The function *vorticity* can be used to calculate vorticity through 4 different numerical schemes. The following commands were used to create figure 2.6:

```
>> w=vorticity(x,y,fu,fv,'circulation');
>> pcolor(x(2:end-1,2:end-1),y(2:end-1,2:end-1),w), shading flat, colorbar
```

To calculate streamlines from the flow field one can use the **Matlab** command *streamline*. This function requires the user to define starting points for each streamline she/he will plot, but this is often a bit tedious

Figure 2.6: Vorticity.

Figure 2.7: Streamlines.

to do manually. **MatPIV** comes with a function named *mstreamline* that finds all the boundaries of the flow and uses these as starting points in *streamline*. The following commands were used to create figure 2.7:

```
>> h=mstreamline(x,y,fu,fv,2);
```

Again the number of vectors in the velocity-field is so large that it is feasible to plot only every second streamline.

2.3 The need-to-know basics

Essentially users end up with just a few lines of code to run **MatPIV** . The following lines assume that we have created both our world-coordinate file (in this example called "worldco.mat") and a mask (in this example called "polymask.mat"). Then the following should suffice:

```
[x,y,u,v,snr]=matpiv('Demo3/mpim1b.bmp','Demo3/mpim1c.bmp',...
[128 128; 64 64; 32 32; 32 32],0.008,1,...
'multin','worldco.mat','polymask.mat');
[su,sv]=snrfilt(x,y,u,v,snr,1.3);
[gu,gv]=globfilt(x,y,su,sv,3);
[lu,lv]=localfilt(gu,gv,2.7,'median',3,'polymask.mat',x,y);
[fu,fv]=naninterp(gu,gv,'linear','polymask.mat',x,y);
```

We note that it is very easy to insert this piece of code into a for loop and perform a large number of calculations. We shall briefly revisit this in section 3.5.

Chapter 3

MatPIV behind the scenes

The present chapter aims to give an overview of the PIV technique and subsequently review the file-structure and the functionality of the different files in the toolbox.

3.1 How does it work?

The fundamentals of Particle Image Velocimetry relies on basic pattern matching. This is a topic touched upon by most introductory books on image processing [see for example Gonzales and Woods, 1992]. The important thing to understand about PIV is that it is nothing but an application of these techniques to experiments in fluid (or solid) mechanics. It is very easy to get confused over the fact that to use PIV you need a grasp of digital cameras, particles, lasers (or other illumination devices) as well as the physics of your experiment. In the present text we aim to emphasize that you should treat the PIV technique as a generic tool to measure your experiments. It is generic in the way that we apply principles of pattern matching (which is a software-issue) and these principles have nothing to do with the way we do our experiment (which depends on hardware-issues). The first step is to add some tracer into our flow field. To measure the motion of this tracer we normally have to illuminate it and film it with a camera. Since most cameras record on a 2-dimensional plane, we usually restrict our illumination to a 2-dimensional plane as well - hereafter known as a "light sheet". To accomplish this we may need to use optics, and to film it we may need some technological insight (like knowing how to use a video camera), but at the end of the day we go through all these steps in order to generate a pattern in our flow field and registering the motion of it. This pattern is subsequently used for pattern matching which will tell us something about displacements. If we then divide the displacement by the time separation between our images (we assume we have taken two images in order to match patterns between them) we end up with a velocity.

3.1.1 Pattern Matching - the principles

Let us start by assuming that we have taken two images (I_1 and I_2), separated by a time distance of $\Delta t = 0.0012$s (like, for example the images in Demo3). We subsequently divide both images into smaller regions, also known as sub-windows, interrogation-windows or interrogation-regions. We then compare each sub-window in the first image with the corresponding sub-window in the second image. We shall hereafter let $I_1^{i,j}$ denote sub-window number i, j in the first image and $I_2^{i,j}$ the corresponding sub-window in the second image.

We now aim to see if we can identify a displacement of the pattern in $I_1^{i,j}$. To do this we can evaluate the squared Euclidean distance between the two sub-windows. This is defined as

$$R_e(s,t) = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} [I_1^{i,j}(m,n) - I_2^{i,j}(m-s,n-t)]^2.$$

This means that, for every possible overlap of the sub-windows, we calculate the sum of the squared difference between them. In other words this means that we are looking for the position where the sub-

windows are the "least unlike". Let us look in a bit more detail to this simple mathematical formula. If we expand the square parentheses on the right hand side we may get

$$R_e(s,t) = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} [I_1^{i,j}(s,t) - I_2^{i,j}(m-s,n-t)]^2$$

$$= \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} I_1^{i,j}(m,n)^2 - 2I_1^{i,j}(s,t) \cdot I_2^{i,j}(m-s,n-t) + I_2^{i,j}(m-s,n-t)^2.$$
(3.1)

We should notice that the first term, $I_1^{i,j}(m,n)^2$, is merely a constant since it does not depend on s and t. The last term, $I_2^{i,j}(m-s,n-t)^2$ is seen to depend on s and t, but we notice that it is just dependent on the second image. So to sum up, only the middle term actually deals with both our images and as a matter of fact this term (without the -2) is usually referred to as cross-correlation (or circular cross-correlation) and defined as

$$R(s,t) = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} I_1^{i,j}(m,n) \cdot I_2^{i,j}(m-s,n-t).$$
(3.2)

The basic assumption here is that the pattern in I_2 is evenly distributed so that the sum of $I_2^{i,j}()^2$ does not change as we vary s and t.

Traditionally in PIV, equation 3.2 has been preferred, and it is also the basis of many of the different algorithms in **MatPIV** ('single','multi' and 'multin' options). The reason for this is primarily that equation 3.2 can be calculated using FFTs (and will therefore execute faster).

The use of equation 3.1.1 in PIV has primarily been advocated by Gui and Merzkirch [1996] and Gui and Merzkirch [2000] and this is implemented in the **MatPIV** option called 'mqd'.

In the field of pattern matching another approach has often been chosen. Considering the fact that the last term in equation 3.1 may be non-constant, many people choose to apply so called normalized correlation, which is defined as

$$R(s,t) = \frac{\sum_{m=0}^{M-1} \sum_{n=0}^{N-1} I_1^{i,j}(s,t) \cdot I_2^{i,j}(m-s,n-t)}{[\sum_{m=0}^{M-1} \sum_{n=0}^{N-1} I_1^{i,j}(m,n)^2 \cdot \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} I_2^{i,j}(m-s,n-t)^2]^{1/2}}.$$
(3.3)

This is the basis in the **MatPIV** option called 'norm'. **MatPIV** here uses a function from the Image Processing Toolbox of **Matlab** called *normxcorr2.m*.

3.2 The core files

3.2.1 matpiv.m

The *matpiv.m* file acts merely as a shell (or batch-file) from which the different calculation files are called. The calling of MatPIV should look like:

```
[x,y,u,v]=matpiv(image1,image2,windowsize,Dt,...
WinOverlap,Method,worldcoordfile,maskfile);
```

Here *image1* and *image2* should be either two matrices containing your preloaded images or two strings containing the names of the images you wish to use.

The parameter *windowsize* should be a number representing the size of the sub-window to be used. However, this depends slightly on the evaluation technique (*Method*) to be applied (see *Method*). *windowsize* is usually set to $16, 32, 64$ or 128.

Dt is the time between the start of exposure of *image1* and *image2*.

WinOverlap is a number between 0 and 1 denoting the overlap of the interrogation regions (sub-windows). Usually this parameter is set to 0.5 or 0.75 which means 50% or 75% overlap of the interrogation windows. We note that *WinOverlap* multiplied by the *windowsize* needs to be an integer number.

Method should be an integer string with one of the following options:

- *'single'* - this performs PIV calculations with a single iteration through the images, using equation 3.2. This equation is also the basis for the method

- *'multi'* - which does PIV with three iterations through the images. This will start off using whatever *windowsize* you specify, but the final iteration will be performed using half this size. An extension of this is called

- *'multin'* - and does PIV with n iterations through the images. In this case *windowsize* needs to be a $n \cdot 2$ sized vector, implying that you have to manually give the size of the sub-window for each iteration. This comes with the added option of using non-quadratic sub-windows. A typical *windowsize*-vector will be something like $[64\ 64;\ 32\ 32;\ 16\ 16;\ 16\ 16]$, which means that we will use 4 iterations starting with 64×64 windows and finishing with 16×16. Note that you can also use something like $[64\ 32;\ 32\ 32;\ 32\ 16 : 32\ 16]$. Additionally we can use the method

- *'mqd'* - which does the PIV calculation in a single pass using equation 3.1.1. Similarly we can use

- *'norm'* - to do the very same thing using equation 3.3.

worldcoordfile should be the name of a file containing the mapping between pixels (camera coordinates) to centimeters (world coordinates). This file is automatically produced when you use the file *definewoco.m*.

maskfile should be the name of a file containing the mask you wish to apply to your images (in other words, containing the region where you do not wish to perform your calculations). This file is produced when you apply the function *mask.m*.

3.2.2 definewoco.m - defining the mapping between pixels and centimeters

The purpose of this file is to calculate the transformation from the local camera coordinates (pixels) to the physical world coordinates in your experiment. Currently only linear mapping is supported. The calling sequence for defining the mapping would then look like:

```
>> definewoco('WorldCoordinateImage.bmp','typeofcoordinate').
```

The world coordinates may be represented either by distinguishable dots or by a regular grid. For the former case one should specify an o for the *typeofcoordinate* input, whilst for the latter one should specify a + sign. If your image is called *WorldCo1.bmp* and contains a regular grid the calling should look like

```
>> definewoco('WorldCo1.bmp','+').
```

This will make a window pop up containing the image. The mouse pointer will change to a circle, and the left mouse button is used to mark the points in the image that mark the World Coordinate system. After you have marked all your points you will be asked to type in the physical coordinates of each of the points marked with the mouse. These coordinates should be on the form [x y], with x being the horizontal and y being the vertical coordinate. Figure 2.2 shows the example image supplied with *MatPIV*. It can be relatively difficult to spot all the points present in the image, The figure shows the points that should be used in your mapping marked out with white circles. This figure also includes the physical coordinates to some of the points as used in that experiment. All numbers are measured in cm.

3.2.3 mask.m - Masking out regions of the flow

The files *mask.m* and *maskpolyg.m* provide a way of masking out a region of your flow after all the calculations have been performed. Ideally this should be done before processing to save calculation time, but this feature will probably take a while before is included in **MatPIV**.

The first step is to create a mask based on the image:

```
[mask,Xverti,Yverti]=mask(Image1,worldcofile);.
```

This will make an image pop up and the user should define the polygon that should be excluded from the measurements. This is done using the left mouse button and the session is ended by pressing the right button. A file called *polymask.mat* will be saved automatically and should be used as input to *maskpolyg.m* when processing the velocity fields:

```
>> [x,y,u,v]=maskpolyg(x,y,u,v,'polymask.mat');.
```

Alternatively one can use the vectors [Xverti,Yverti] as input directly:

```
>> [x,y,u,v]=maskpolyg(x,y,u,v,[Xverti,Yverti]);.
```

3.3   Filters in MatPIV

After having performed the calculations it is usually necessary to filter the velocity data in order to remove so called spurious vectors (aka outliers). This can be done in **MatPIV** using four different filters:

1. Signal-To-Noise ratio filter,

2. Peak height filter,

3. global filter and

4. local filter.

The two latter of these files include various methods, global or local, for removing outliers. In this section we will take a look at how to use the velocity filters in **MatPIV**. We note that all the filters will replace the spurious vectors with NaNs (Not A Number).

3.3.1   Signal-to-Noise ratio

The output from **MatPIV** can, apart from the velocities, also include the Signal-To-Noise ratio and the peak-height for use in validation of the vector field. The file *snrfilt.m* is used to validate the vector field with regards to the Signal-To-Noise ratio output from *matpiv*. The calling sequence looks like:

```
[su,sv]=snrfilt(x,y,u,v,snr,threshold,actions);
```

Here the final input *actions* can be ommited (typically in batch processing), or specified as 'loop'. The 'loop' option will allow the user to change the threshold value interactively and view the effect the changes has on the vector field. All the "invalid" velocity vectors (with a SnR lower than the threshold) are replaced with NaN (Not A Number). Keane and Adrian Keane and Adrian [1992] suggest that a threshold value of about 1.3 is appropriate, so that normally something like

```
>> [su,sv]=snrfilt(x,y,u,v,snr,1.3);
```

will give good results, but this will obviously depend on the quality of your images.

3.3.2   Global histogram operator

Globfilt is a so called global histogram operator that will remove vectors that are significantly larger or smaller than a majority of the vectors.

There are two basically different methods used in *Globfilt*, The first uses a graphical input to specify the acceptance interval of velocity vectors. These are plotted in the (u,v) plane and the user should specify 4 points, using the left mouse button, that together form a 4 sided polygonal region of acceptance. Use 'manual' as input parameter for this option. Alternatively one can make an acceptance interval, either based on the standard deviations (x and y) of the measurement ensemble or simply by specifying the limits directly. This can be done in three different ways, namely 1) by specifying a factor (a number), 2) by specifying 'loop' or 3) by specifying a vector with the upper and lower velocity limits. In the former case *Globfilt* uses the mean of the velocities plus/minus the number times the standard deviation as the limits

for the acceptance area. In the second case *Globfilt* loops and lets the user interactively set the factor. This option often performs well if the vector field is not heavily contaminated with outliers. The third option is used by specifying an input vector [Umin Umax Vmin Vmax] which defines the upper and lower limits for the velocities.

```
[gu,gv]=globfilt(x,y,su,sv,actions);
```

where *x* and *y* are the coordinate matrices and *su* and *sv* the filtered velocity matrices output from the snr-filter.

Actions can be (as mentioned above)

1. the string 'manual',

2. a scalar (threshold) value,

3. the string 'loop' or

4. a vector.

The former option let's the user specify an acceptance region graphically using the left mouse button. The second option, a scalar, is used to create an acceptance region from the formula $\mathbf{U_{min/max}} = \text{Factor} * \mathbf{std(U)}$. This option is nice for batch processing. Alternatively one might use the 'loop' option. This option can additionally take an initial scalar as mentioned above and will loop to let the user decide the acceptance region based on this scalar/threshold. Finally one can simply choose to specify upper and lower limits on the velocities. This should be done in a vector [Umin Umax Vmin Vmax].

In real applications the following will usually produce good results:

```
>> [gu,gv]=globfilt(x,y,su,sv,4);
```

3.3.3 Local filter

The *localfilt*-file incorporates the two different filters, namely a median filter and a mean filter. They filter velocities based on the squared difference between individual velocity vectors and the median or the mean of their surrounding neighbors.

This filter is called with the following parameters:

```
[mu,mv]=localfilt(u,v,threshold,method,kernelsize,mask,x,y);
```

Typically we consider a region *kernelsize * kernelsize* large and compare the vector in the middle with the remaining vectors using one of the two *methods* 'median' or 'mean'. The threshold determines wich vectors are thrown out. In the following example we'll use a 3*3 kernel and a threshold of 2.5:

```
>> [lu,lv]=localfilt(gu,gv,2.5,'median',3);
```

In a vector field, let us consider vector number i, j and compare it to the 8 vectors surrounding it (3*3 kernelsize) using the median-option. Then the vector is rejected if it is larger than the median of all the 9 vectors plus the threshold times the standard deviation of all the vectors, or if it is smaller than the median minus the threshold times the standard deviation. Mathematically we can say that a vector is considered an outlier if

$$\mathbf{U}_{i,j} \gtrless \text{median}(U_{i-1:i+1,j-1:j+1}) \pm \text{threshold} \cdot \text{std}(U_{i-1:i+1,j-1:j+1}).$$

This filter implies that any vector can not be "too different" from its neighbors. The value of the threshold will determine exactly *how* different. A value between $1.7 and 3$ will usually be sufficient but users should keep in mind that large values here will result in fewer outliers than a small number.

The kernelsize may be chosen as any odd number, but practically speaking a value of 3 (meaning a 3*3 kernel) or 5 will do just fine for most users.

Finally the input to localfilt can also include the name of the mask-file in order to save some time. If this is not specified, localfilt will loop through all the elements in the velocity matrices, even if they have already been classified as outliers by earlier filters. In this case it is vital that also the coordinate matrices, *x* and *y*, are included. Here's an example:

```
>> [lu,lv]=localfilt(gu,gv,2.5,'median',3,'polymask.mat',x,y);
```

The median of the vectors can easily be replaced with the mean value, although the former is the default method since it is more robust to other outliers in the neighborhood (see Westerweel et al. [1997]).

3.3.4   Interpolate outliers

The file *naninterp.m* interpolates NaN's in a vectorfield using two slightly different methods. The calling should look like

```
[fu,fv]=naninterp(lu,lv,method,mask,x,y);,
```

where method should be 'linear' or 'weighted' and defaults to 'linear' if not specified. This option sorts all spurious vectors based on the number of spurious neighbors to a point and starts by interpolating the vector that have as few neighboring outliers as possible, looping until no NaN's are left. The 'weighted' method uses the file FILLMISS.M to interpolate the outliers (see 'help fillmiss' for some info).

Additionally one can apply a mask to avoid interpolating vectors in a part of the measurement area.

```
>> [u,v]=naninterp(u,v,'polymask.mat',x,y);,
```

will replace NaN's but leave out areas that have been masked out (using MASK.M) Using the MASK option requires the x and y matrices input along with the velocities, u and v.

3.4   Integral and differential quantities

3.4.1   Streamlines

To calculate streamlines from your flow please refer to the MATLAB file streamline.m (type *help streamline* at the command prompt). This file calculates streamlines from a given set of starting points, and to ease the plotting **MatPIV** comes with a file calles mstreamline.m which will create a set of starting points on all the edges of the flow field.

The file streaml.m, originally included with **MatPIV** has been discontinued and will no longer be updated.

3.4.2   Vorticity

We start by exploiting the possibility of extracting the vorticity from our 2-D measurement. Vorticity can be estimated by calculating

$$\omega = \frac{\partial V}{\partial X} - \frac{\partial U}{\partial Y}.$$

Several numerical schemes exist for performing this calculation, and three different methods have been implemented in **MatPIV**. The first is estimation by using Stokes theorem:

$$\omega_{i,j} = \frac{1}{8 * \Delta X \Delta Y}[\Delta X(U_{i-1,j-1} + 2U_{i,j-1} + U_{i+1,j-1})$$

$$+\Delta Y(V_{i+1,j-1} + 2V_{i+1,j} + V_{i+1,j+1})$$

$$-\Delta X(U_{i+1,j+1} + 2U_{i,j+1} + U_{i-1,j+1})$$

$$-\Delta Y(V_{i-1,j+1} + 2V_{i-1,j} + V_{i-1,j-1})].$$

This approach integrates to find the circulation around a point. Alternatively we can use a standard differentiation scheme, such as forward or centered differences. **MatPIV** uses two different and more accurate differential operators, namely least squares and Richardson extrapolation. With the former of these schemes the vorticity can be estimated by

$$\omega_{i,j} = \frac{1}{10\Delta X}(2v_{i+2,j} + v_{i+1,j} - v_{i-1,j} - 2v_{i-2,j})$$

$$-\frac{1}{10\Delta Y}(2u_{i,j+2} + u_{i,j+1} - u_{i,j-1} - 2u_{i,j-2}),$$

while the latter uses

$$\omega_{i,j} = \frac{1}{12\Delta X}(-v_{i+2,j} + 8v_{i+1,j} - 8v_{i-1,j} + v_{i-2,j})$$

$$-\frac{1}{12\Delta Y}(-u_{i,j+2} + 8u_{i,j+1} - 8u_{i,j-1} + u_{i,j-2}).$$

The major difference between the two last operators is that the Richardson extrapolation is designed to produce a smaller truncation error, while the least squares operator reduces the effect of fluctuations. The latter reason is why the least squares operator is often used with PIV measurements and is chosen as the default calculation method in **MatPIV**.

The file *vorticity.m* calculates vorticity. Calling should look like:

```
[vorticity]=vorticity(x,y,u,v,method).
```

Method should be one of 'centered', 'circulation', 'richardson' or 'leastsq', where the latter is the default. The former of the methods is just an overloaded call to the MATLAB file curl.m (type *help curl* at the command prompt). This file uses a centered differences approach and will perform well if the velocities are smooth.

If no output argument is included, a figure window will appear showing the result.

3.4.3   Other files included

vekplot2.m

File written by Per-Olav Rusaas (peolav@math.uio.no). Essentially equivalent to QUIVER but with the difference that scaling is NOT optional. Therefore the scaling is always known and it is easy to plot two vector fields on top of each other with identical scaling.

mginput.m

File written by Per-Olav Rusaas (peolav@math.uio.no). A small change to the original GINPUT to use a circle as the mouse cursor. Used with the *definewoco.m* file.

xcorrf2.m

File written by R. Johnson (no e-mail address available at the time of writing). The file can be found at the Mathworks web-site: *http://www.mathworks.com/*.

fillmiss.m

File written by Kirill K. Pankratov, kirill@plume.mit.edu (1994). File found at the Mathworks web-site. File performs a linear interpolation of NaN's in a matrix. The help text in the file says (quote): *"Missing values are calculated as weighted sum of linear interpolations from nearest available points. Altogether 5 estimates from column-wise and 5 for row-wise 1-d linear interpolation are calculated. Weights are such that for the best case (isolated missing points away from the boundary) the interpolation is equvalent to average of 4-point Lagrangian polynomial interpolations from nearest points in a row and a column."*

mnanmedian.m and mnanmean.m

Code suggested by John Peter Acklam, jacklam@math.uio.no at the comp.soft-sys.matlab newsgroup, 1999. Finds the median/mean of a vector or matrix with NaN-elements. These files are replacements to the original nanmedian and nanmean files which come with the **MATLAB** statistical toolbox.

3.5   Batch-processing with MatPIV

One of the advantages of using **MATLAB** is, besides the obvious platform independency, the ability to batch-process vast amounts of data by writing script-files. In the previous chapter we used a few "manual" options during filtering and this is something that we would prefer to avoid when we work with larger amounts of measurement data. In this case **MatPIV** has the ability to work as a background process performing all the measurements. For instance the calculations performed in the previous chapter could have been accomplished by writing and executing a script file containing the lines:

```
[x,y,u,v,snr]=matpiv('im00.bmp','im04.bmp',...
64,0.04,0.5, 'mqd','worldco1.mat');
[su,sv]=snrfilt(x,y,u,v,snr,1.1);
[gu,gv]=globfilt(x,y,su,sv,3);
[lu,lv]=localfilt(gu,gv,2,'interp');
[fu,fv]=naninterp(lu,lv,'linear');
```

In this case we assume that the world coordinate file already exists, but we have not used a mask. We have also applied the MQD-method for pattern matching, just to emphasize that it is possible.

Similarly it is rather easy to write *for*-loops in **MATLAB** to repeat this operation several times, replacing the input images each time. **MatPIV** is supplied with an example of such a file, which is called *runfile.m* and is included in the distribution. Typically such a loop can be written in the following way:

```
%%%%%%%%%%%%%%%%%%%% Settings
T=0.04;        % Time separation between base and cross images.
met='multin'; % Use interrogation window offset
wins=[64 64;64 64;32 32;32 32]; %windowsizes to use in the iterations
woco='worldco1.mat'; % World coordinate file. This may be changed
                     % within the loop as well, e.g. in the cases
                     % where there are two cameras
snrtrld=1.2;  % threshold for use with snr-filtering
globtrld=3;   % threshold for use with globalfiltering
loctrld=1.7;  % threshold for use with local filtering
med='median'; % Use median filtering in localfilt
int='linear'; % interpolate outliers using linear interpolation
maskfile='polymask.mat'; % name of file containing the pre-defined mask
M=length(dir('*.bmp'));  % check how many images are present in the directory
%%%%%%%%%%%%%%%%%%%%%
for i=1:M-1
  im1=['myfilename',num2str(i),'.bmp'];
  im2=['myfilename',num2str(i+1),'.bmp'];
  % PIV calculation
  [x,y,u,v,snr]=matpiv(im1,im2,wins,T,0.5,met,woco,maskfile);'])
  % SnR filter:
  [su,sv]=snrfilt(x,y,u,v,snr,snrtrld);
  % global filter:
  [gu,gv]=globfilt(x,y,su,sv,globtrld);
  % local median filter:
  [lu,lv]=localfilt(gu,gv,loctrld,med,x,y);'])
  % interpolate outliers
  [fu,fv]=naninterp(lu,lv,int,maskfile,x,y);'])

  save(['vel_field',num2str(i),'.mat'],'u','v','snr','fu','fv');

  if i==1
     save coordinates.mat x y
```

```
      end
end
```

   In this way we will automatically save one file for each velocity field. This file will contain both the unfiltered and the filtered velocities, since this will enable us to filter the data at a later stage with different parameters (should that be necessary). If we assume that the number of images (M) is large it should be clear that this way of processing has it's advantages. The coordinate arrays, x and y, will be saved in a separate file, coordinates.mat.

   If we, on the other hand, are working with AVI-movies, this example shows how our script could work:

```
%%%%%%%%%%%%%%%%%%%% Settings
T=0.04;        % Time separation between base and cross images.
met='multin'; % Use interrogation window offset
wins=[64 64;64 64;32 32;32 32]; %windowsizes to use in the iterations
woco='worldco1.mat'; % World coordinate file. This may be changed
                     % within the loop as well, e.g. in the cases
                     % where there are two cameras
snrtrld=1.2;  % threshold for use with snr-filtering
globtrld=3;   % threshold for use with globalfiltering
loctrld=1.7;  % threshold for use with local filtering
med='median'; % Use median filtering in localfilt
int='linear'; % interpolate outliers using linear interpolation
maskfile='polymask.mat'; % name of file containing the pre-defined mask
A=aviread('myfile.avi')); % load the movie. If the movie is large
                          % this should be done inside the loop
%%%%%%%%%%%%%%%%%%%%
for i=1:length(A)-1
  % PIV calculation
  [x,y,u,v,snr]=matpiv(A(i).cdata,A(i+1).cdata,wins,T,0.5,met,woco,maskfile);'])
  % SnR filter:
  [su,sv]=snrfilt(x,y,u,v,snr,snrtrld);
  % global filter:
  [gu,gv]=globfilt(x,y,su,sv,globtrld);
  % local median filter:
  [lu,lv]=localfilt(gu,gv,loctrld,med,x,y);'])
  % interpolate outliers
  [fu,fv]=naninterp(lu,lv,int,maskfile,x,y);'])

  save(['vel_field',num2str(i),'.mat'],'u','v','snr','fu','fv');

  if i==1
     save coordinates.mat x y
  end
end
```

3.6 Using a parameter-file as input

Finally we shall mention that it is also possible to supply an m-file containing your parameters as input to matpiv. The file *parameters_example.m* gives an example of how this may be achieved with the images in the Demo 1 directory. The user would then execute the following:

```
>> [x,y,u,v,snr]=matpiv('parameters_example.m');
```

   Here the file *parameters_example.m* is just a list of the parameters we would normally write on the command line:

```
% Parameter-example file
im1= 'Demo1/im00.bmp'; %The "base" image
im2= 'Demo1/im04.bmp'; %The "cross" image
Dt=0.04; %Time between images
winsize=[64 64;64 64;32 32;32 32]; %interrogation region size
overlap=0.5; % Overlap of interrogation regions
method='multin'; %Method for interrogation (i.e. multiple passes)
wocofile='Demo1/worldco.mat'; %file containting the mapping from image to
                        %world coordinates
msk='Demo1/polymask.mat'; %file defining the regions to mask from the
                        %flow.
```

Bibliography

Rafael C. Gonzales and Richard E. Woods. *Digital Image Processing*. Addison-Wesley, 1992.

J. Grue, A. Jensen, P-O. Rusås, and J.K. Sveen. Properties of large amplitude internal waves. *J. Fluid Mech.*, 380:257–278, 1999.

L. Gui and W. Merzkirch. Generating arbitrarily sized interrogation windows for correlation-based analysis of particle image velocimetry recordings. *Exp. Fluids*, 21:465–468, 1996.

L. Gui and W. Merzkirch. A comparative study of the mqd method and several correlation-based piv evaluation algorithms. *Exp. Fluids*, 28:36–44, 2000.

Atle Jensen, J. Kristian Sveen, John Grue, Jean-Baptiste Richon, and Callum Gray. Accelerations in water waves by extended particle image velocimetry. *Experiments in fluids*, 30:500–510, 2001.

Richard D. Keane and Ronald J. Adrian. Theory of cross-correlation analysis of piv images. *Applied Scientific Research*, 49:191–215, 1992.

M. Raffel, C. E. Willert, and J. Kompenhans. *Particle Image Velocimetry, A Practical Guide*. Springer Verlag, first edition, 1998.

J. K. Sveen and E. A. Cowen. Quantitative imaging techniques and their application to wavy flow. In J. Grue, P. L. F. Liu, and G. K. Pedersen, editors, *PIV and Water Waves*. World Scientific, 2004.

J. Westerweel. *Digital Particle Image Velocimetry- Theory and Application*. PhD thesis, Delft University of Technology, The Netherlands, 1993.

J. Westerweel, D. Dabiri, and M. Gharib. The effect of a discrete window offset on the accuracy of cross-correlation analysis of digital piv recordings. *Exp. Fluids*, 23:20–28, 1997.